

# Decontaminated by Whose Standard? Label-Free Measurement of Decontamination Filter Error

Anonymous ACL submission

## Abstract

Open post-training projects routinely state that their training data has been decontaminated against public benchmarks, and readers treat that statement as meaning roughly the same thing across projects. Reimplementing five such filters from their released source and running all of them over one corpus against one benchmark suite, we find they flag between 0% and 27.5% of the same items. Disagreement alone, however, invites the reply that different tools simply sit at different thresholds.

Our main contribution settles that. Benchmark items published *after* a corpus was assembled cannot appear in it, so any flag against them is a false positive by construction, identifiable without annotation. Using competition sets held after the corpus cutoff as a negative control, we find the two filters that detect anything are wrong about more than a tenth of what they report: 12.9% and 11.1% false positives, replicated at 6.2% and 6.9% on a second, independent corpus. That is not a threshold preference but a measured error rate, and it is stable in proportion across a  $31\times$  range of corpus size. We verify the control assumption directly rather than assuming it, and trace every error to the stock phrasing of competition mathematics rather than to leaked content.

Because the control is label-free, it also tunes. Applied as a sweep it shows that  $n$ -gram size trades sensitivity without buying precision, and that the best-performing method in our comparison ships a default threshold that disables it.

## 1 Introduction

A benchmark score is evidence of capability only if the benchmark was held out. Open post-training projects generally assert that it

was: model cards and papers state that training data was decontaminated against the suites on which results are reported. Readers, and downstream leaderboards, treat those assertions as comparable.

They are comparable only if the underlying procedures agree. This paper asks whether they do, and finds that they do not, by a wide margin.

We take five decontamination filters that ship with real released datasets, reimplement each from its own source, and apply all five to a single corpus against a single benchmark suite. The spread is not subtle. On 100,000 training documents and 502 benchmark items, the filters flag 0, 0, 6, 127 and 138 items respectively. Whether a benchmark problem counts as contaminated is, in practice, a function of which tool was run.

That much is only an observation, and it has an easy dismissal: different tools are tuned differently, and a spread of thresholds is what one would expect. Establishing that filters disagree does not establish that any of them is *wrong*.

Answering that has previously required human adjudication, which is expensive and unreliable at exactly the boundary cases that decide it. We propose a cheaper and more decisive test. Competition mathematics produces new benchmarks continuously and with known dates. An item published after a corpus was assembled cannot be present in it. Any filter that flags such an item has made an error that requires no judgement to identify, and the dismissal does not survive it, because a threshold choice does not explain why half of what a filter reports is wrong.

Using this construction we contribute:

- A label-free method for measuring decontamination false-positive rates, using post-

084  
085  
086  
  
087  
088  
089  
090  
091  
  
092  
093  
094  
095  
  
096  
097  
098  
099  
  
100  
101  
102  
103  
104  
105  
  
106  
107  
108  
109  
110  
  
111  
  
112  
113  
114  
115  
116  
117  
118  
119  
120  
121  
122  
123  
124  
125  
126  
127  
128  
129  
130

cutoff benchmarks as a negative control (§4), with the control assumption tested rather than assumed (§4.1).

- Measured error rates of 12.9% and 11.1% for the two filters that detect anything, replicated at 6.2% and 6.9% on an independent corpus (§9), and stable in proportion across a  $31\times$  range of corpus size (§8).
- A mechanistic account of those errors: all of them arise from a small set of conventional competition-mathematics phrases (§10).
- A direct comparison of five published decontamination filters, each run from its released implementation on identical inputs (§6).
- A tuning procedure built on the control set, showing that  $n$ -gram size is a sensitivity parameter rather than a precision parameter, and that fuzzy whole-string matching dominates  $n$ -gram matching at every operating point we measured (§11).

We emphasise at the outset that this is not a finding that any project was careless. Every filter we study does what its authors implemented. The problem is that one word describes procedures measuring different things.

**2 Related Work**

Work on benchmark contamination divides into detection from the model side and inspection from the data side. Model-side methods ask whether a trained model shows evidence of having seen an evaluation item, through membership inference, loss-based tests, or exchangeability arguments (Oren et al., 2024). A systematic review of 55 such studies finds that no method is reliable across contamination regimes (Nourbakhsh et al., 2026), and Zarzecki et al. (2026) report that across 335 evaluations of three leading paradigms only 60% yield correct outcomes, concluding explicitly that statistical detection cannot yet replace transparent data provenance. Shihab et al. (2026) give an information-theoretic account of why, showing that detectability scales with prevalence and sample size such that audit-scale samples are underpowered by one

to two orders of magnitude. Contamination detection in reasoning models is separately fragile (Wang et al., 2026), and detectors built on LLM judges inherit contamination problems of their own (Li et al., 2026a, 2025).

This paper works on the provenance side that such conclusions point toward: the training corpus is public, so overlap can be inspected directly. Prior data-side work has generally audited one corpus with one detector and reported a contamination rate (Sainz et al., 2023), and paraphrase-based evasion is known to defeat surface matching (Yang et al., 2023). Our question is prior to those. Before asking how contaminated a corpus is, we ask whether the quantity is well defined, given that the procedures in use disagree. Recent work pursuing statistical guarantees for detection (Zhang et al., 2026) is complementary: it assumes the target quantity is fixed, which is what we test.

A parallel literature examines the reliability of audit pipelines themselves. Li et al. (2026b) give a five-class taxonomy of pipeline failures in perturbation-based validity audits, and Zarzecki et al. (2026) study distribution shift and scale as limits on model-side detection. Both concern audits of models. To our knowledge no prior work runs multiple published decontamination *implementations* over a common corpus and reports the divergence in what they flag.

The artifacts we study are s1 (Muennighoff et al., 2025), Tulu 3 (Lambert et al., 2024), OpenThoughts and OpenThoughts3 (Guha et al., 2025), and Olmo 3 (Team Olmo, 2025), whose data and decontamination code are all public. Our benchmarks come from OlymMATH (Sun et al., 2025) and MathArena (Dekoninck et al., 2026).

**3 The Filters**

We reimplemented each filter from its released source, or compiled and ran the released binary where one exists. Table 1 summarises what each does, which in two cases differs from what its paper describes.

Three observations follow from reading the sources.

**The filters do not implement the same kind of test.** Three are single-stage: a match on any shared  $n$ -gram, or a whole-

131  
132  
133  
134  
135  
136  
137  
138  
139  
140  
141  
142  
143  
144  
145  
146  
147  
148  
149  
150  
151  
152  
153  
154  
155  
156  
157  
158  
159  
160  
161  
162  
163  
164  
165  
166  
167  
168  
169  
170  
171  
172  
173  
174  
175  
176  
177  
178  
179  
180

| Filter            | Method                               | Tokenizer              | Params from      |
|-------------------|--------------------------------------|------------------------|------------------|
| s1                | any shared 8-gram                    | whitespace, lowercased | code             |
| Tulu 3            | coverage fraction over Elasticsearch | spaCy / ES analyzer    | paper prose only |
| OpenThoughts-114K | <code>fuzz.ratio</code> $\geq 95$    | none (raw strings)     | code             |
| OpenThoughts3     | any shared 13-gram                   | Qwen/Qwen2-7B-Instruct | code             |
| decon (Olmo 3)    | 5-gram detect, then cluster scoring  | cl100k                 | code             |

Table 1: The five filters as implemented rather than as described. Two points deserve emphasis: OpenThoughts-114K performs no  $n$ -gram matching at all, and Tulu 3’s two governing parameters have no defaults anywhere in its released code.

string similarity above a cutoff. **decon** is two-stage, generating candidates by 5-gram overlap and accepting them only if a weighted score exceeding 0.8 survives cluster expansion. Tulu 3 computes a coverage fraction over documents retrieved by Elasticsearch, so its recall is bounded by retrieval.

**The tokenizer is where single-stage filters diverge most.** s1 splits on whitespace after lowercasing; OpenThoughts3 hardcodes a model tokenizer. Eight whitespace tokens and eight model tokens describe very different amounts of text, so the same nominal  $n$  means different things.

**In two cases the released code and the written description do not line up.** These are observations about artifacts, not about anyone’s conduct, and in both cases there is a benign reading.

The decontamination code released with OpenThoughts-114K contains no  $n$ -gram logic at all: it calls `process.extract` with `fuzz.ratio` at a cutoff of 95 on raw strings. A 13-gram path does exist in the same repository, but under OpenThoughts3, a later and separate release. A reader who takes “13-gram” as describing the 114K artifact would be mistaken, though the description is accurate for the artifact it belongs to. Distinguishing the two requires reading the repository.

Tulu 3’s `--ngram_size` and `--match_threshold` are declared without defaults and fall through as `None`, so the “8-gram” and “> 50% coverage” that its paper states are not recoverable from the code. The values were presumably supplied at run time by a driver that was not released. Either way a third party cannot reproduce the configuration, so we mark Tulu 3’s parameters as prose-derived throughout and treat its

numbers accordingly.

Neither observation implies the filters were applied carelessly. Both illustrate the paper’s point: what “decontaminated” denotes is fixed by an implementation, and an implementation is only as recoverable as what was released alongside it.

## 4 Label-Free Evaluation via Temporal Controls

Determining whether a flagged item is genuinely contaminated normally requires human adjudication, which is costly and unreliable at the boundary cases that matter most.

Competition mathematics offers an alternative. New problem sets are released continuously, at publicly known dates. For a corpus assembled at time  $t$ , any benchmark published after  $t$  cannot be contained in it, whatever a filter reports. Flags against such items are false positives by construction: not estimates, not adjudications, but logical impossibilities.

This yields a two-group design. A **measurement set** of benchmarks published before the corpus cutoff, where flags may be correct or incorrect, and a **control set** published after it, where every flag is wrong. The control set gives a false-positive rate directly, with no annotation, and can be recomputed for any corpus whose assembly date is known.

The design has one requirement: the cutoff must be established independently of the filters being evaluated. We use the corpus’s public release date as an upper bound, and exclude benchmarks falling near it.

### 4.1 Testing the assumption

That upper bound is load-bearing, and Dolci is a remix of separately curated corpora, so in principle a component could postdate the release of the whole. We therefore test it rather

than assert it. A benchmark item genuinely present in a corpus shares long spans with some training document; a coincidental match shares only short ones. For every benchmark item we computed the longest  $n$  for which it shares an  $n$ -gram with any training document.

| Longest match | Measure<br>(502) | Control<br>(280) |
|---------------|------------------|------------------|
| no match      | 375              | 244              |
| $\geq 8$      | 91               | 25               |
| $\geq 10$     | 29               | 9                |
| $\geq 13$     | 7                | 1                |
| $\geq 15$     | 0                | 1                |

Table 2: Longest word  $n$ -gram each benchmark item shares with the corpus.

No control item matches at a length that would indicate presence, and inspection of the five longest control spans finds no problem content at all. The longest, at  $n=15$ , is “*where  $m$  and  $n$  are relatively prime positive integers, find  $m + n$* ”, the conventional closing of an AIME problem. The others are “*is not divisible by the square of any prime*” and “*reads the same from left to right as from right*”. These are genre conventions, not leaked text.

The longest single shared span anywhere in this study therefore belongs to an item that provably postdates the corpus, and consists entirely of boilerplate. That is both the assumption discharged and the clearest single illustration of the mechanism analysed in §10.

## 5 Experimental Setup

**Corpus.** allenai/Dolci-Think-SFT-32B, the post-training data for Olmo 3 (2,253,684 examples). Olmo 3 was announced 20 November 2025, so we treat the corpus as closed by roughly October 2025. We use a 100,000-document sample for the main results and up to 390,000 for the scaling analysis. Filters read the first user turn of each example, the field their own implementations target.

Dolci is a remix of separately curated corpora. In our sample the largest slices are OpenThoughts3 Math (34,835), Python Algorithms (21,090), Persona Precise IF (10,089), Nemotron Code (5,232) and SYNTHETIC-2 (4,914). Each was decontaminated, if at all, by whoever assembled it.

| Filter                | Measure<br>(502) | Control<br>(280) |
|-----------------------|------------------|------------------|
| decon                 | 0 (0.0%)         | 0 (0.0%)         |
| fuzzy, cutoff 95      | 0 (0.0%)         | 0 (0.0%)         |
| Tulu 3 coverage       | 6 (1.2%)         | 0 (0.0%)         |
| s1 8-gram             | 127 (25.3%)      | 36 (12.9%)       |
| OpenThoughts3 13-gram | 138 (27.5%)      | 31 (11.1%)       |

Table 3: Benchmark items flagged, 100,000 Dolci documents. Every control-set flag is a false positive by construction.

**Measurement set (502 items).** Olym-MATH (200) plus ten MathArena competitions held on or before September 2025: Apex 2025, the Apex shortlist, SMT 2025, CMIMC 2025, BRUMO 2025, AIME 2025, HMMT February 2025, both AIME 2024 papers, and HMMT February 2024.

**Control set (280 items).** Eight MathArena sets from January to August 2026: AIME 2026, HMMT February 2026, USAMO 2026, three ArXivMath sets and two BrokenArXiv sets. All postdate the corpus by at least three months.

## 6 Filters Disagree

Table 3 gives the central result. The filters separate into two groups with nothing between them. Three flag almost nothing and are never wrong. Two flag a quarter of the benchmark and are wrong about more than a tenth of it. There is no shared standard for where on this range the word “decontaminated” should sit.

Agreement between them is close to absent. Of the 201 items (40.0% of the benchmark) flagged by at least one filter, 133 (66.2%; 95% CI [59.6, 72.6] by bootstrap over benchmark items,  $B=10,000$ ) are flagged by *exactly one*. Two items reach three of five, and none reaches four. The highest pairwise Jaccard overlap is 0.318, between the two  $n$ -gram filters; every other pair falls below 0.05.

decon’s zero is robust rather than marginal: it returns zero at its default score threshold of 0.8, zero at 0.6, a single contaminated document at 0.4, and 646 only at 0.2, a quarter of the default.

**A caution about partial reimplementations.** Our first attempt at decon implemented only its candidate-generation stage, a

5-gram match over c1100k tokens. That version flagged 199 of 200 items. The real binary flags none. Reporting the approximation would have inverted the result and misrepresented the tool. A filter is its whole pipeline, and two-stage filters cannot be approximated by their first stage.

## 7 Measured False-Positive Rates

Both  $n$ -gram filters flag items that cannot be in the corpus: 12.9% of the control set for s1 (95% CI [9.3,16.8]) and 11.1% for OpenThoughts3 ([7.5,15.0]), against measurement rates of 25.3% ([21.5,29.1]) and 27.5% ([23.5,31.7]). Intervals are bootstrap over benchmark items,  $B=10,000$ ; the control set has 280 items, so these carry roughly four points of uncertainty and both nonetheless exclude zero by a wide margin.

Subtracting the control rate from the measurement rate leaves roughly 12.4% and 16.4% as estimates of genuine detection, meaning about half of s1’s reported signal is noise. The three conservative filters have no false positives and almost no true positives.

## 8 The Error Rate Is Intrinsic

A single-corpus measurement invites the objection that it may not survive scaling. It does not need to: Figure 1 shows the error fraction is invariant. Across seven corpus sizes from 12,500 to 390,000 documents, absolute rates roughly double, for false positives as much as for detections, while the *proportion* of error stays flat. s1’s false-positive fraction runs 0.49, 0.52, 0.61, 0.51, 0.53, 0.53, 0.54; OpenThoughts3’s runs 0.41, 0.36, 0.37, 0.40, 0.40, 0.43, 0.46, drifting upward, so that filter becomes relatively worse as the corpus grows.

The corrected signal does grow, from 7.0% to 12.8% for s1 and 7.6% to 19.7% for OpenThoughts3, so real overlap is found beneath the noise. It is outnumbered at close to a fixed ratio, and no choice of corpus size changes that.

`fuzz.ratio` returns zero at every scale up to 250,000 and one item at 390,000; `decon` returns zero throughout. The architectural disagreement between whole-string and  $n$ -gram matching is therefore scale-invariant, and is the more portable of our two findings.

## 9 The Errors Are Not Specific to One Corpus

To test whether the false-positive rate is a property of these filters or of Dolci, we repeated the measurement on an independent corpus: the Tülu 3 SFT mixture’s mathematics subset, released November 2024. The date split moves with the corpus. Every 2025 competition now postdates it, so the control side grows to 722 items and only the three 2024 competitions remain as a measurement set.

| Filter                | Detect<br>(60) | False pos.<br>(722) |
|-----------------------|----------------|---------------------|
| s1 8-gram             | 2 (3.3%)       | 45 (6.2%)           |
| OpenThoughts3 13-gram | 1 (1.7%)       | 50 (6.9%)           |
| fuzzy, cutoff 95      | 0 (0.0%)       | 0 (0.0%)            |
| Tülu 3 coverage       | 0 (0.0%)       | 0 (0.0%)            |

Table 4: The same filters on the Tülu 3 SFT math subset, 100,000 documents.

Both  $n$ -gram filters again flag items that cannot be present, at 6.2% and 6.9%. The phenomenon is therefore a property of the method rather than of Dolci. Its magnitude is not: the rate here is roughly half what it is on Dolci, so a false-positive rate must be measured per corpus and cannot be carried over.

We deliberately draw no conclusion about detection on this corpus. Sixty measurement items give an estimate with a 95% interval of roughly [0.4,11.5] percent, which cannot be compared meaningfully against the false-positive rate. The control side is what this corpus is useful for, and it is the larger of the two.

Incidentally, Tülu 3’s own filter flags nothing in Tülu 3’s own data, which is what one would expect if it was in fact run.

## 10 Why the Errors Happen

All 36 of s1’s false positives trace to 73 distinct 8-grams, and none is a fragment of a leaked problem. They are the conventional phrasing of competition mathematics: *find the sum of all possible values of* (166 documents), *let  $a$  and  $b$  be positive integers such that* (67), *is not divisible by the square of any* (67), *where  $m$  and  $n$  are relatively prime positive* (8 documents, matching 6 distinct control items).

AIME problems conventionally close by asking for  $m + n$  where  $m$  and  $n$  are relatively




Figure 1: Across a  $31\times$  range of corpus size, both the flag rate and the false-positive rate roughly double (a), while the fraction of each filter’s output that is provably wrong stays flat (b). The error is a property of the method, not of how much data it is given.

prime; “not divisible by the square of any prime” is standard framing for a radical-form answer. These conventions are shared by every problem of the type, including problems written months after the corpus was assembled.

This explains the precision ceiling of §11. Competition mathematics has a fixed inventory of set phrases eight or more words long, so an eight-word window can be filled entirely by convention, carrying no problem-specific content. Increasing  $n$  does not escape the problem, because the phrases are themselves long.

## 11 Tuning Without Labels

Because control-set flags are provably wrong, sweeping a filter’s parameters over both groups yields a precision/recall curve with no annotation anywhere (Figure 2).

**$n$ -gram size is a sensitivity knob, not a precision knob.** For whitespace tokens the error fraction is pinned near 0.51 from  $n=8$  to  $n=13$  while detection falls from 25.3% to 1.4%: raising  $n$  buys no precision, only insensitivity. At  $n=15$  detection reaches zero while false positives persist, so everything found is wrong. Model-token  $n$ -grams do improve with  $n$ , from an error fraction of 1.00 at  $n=5$  to 0.40 at  $n=13$  and 0.30 at  $n=25$ , but reaching a fraction below 0.1 costs all but 0.4% of detection.

Notably, s1 (whitespace,  $n=8$ : 25.3% detection, 12.9% error) and OpenThoughts3 (Qwen,

$n=13$ : 27.5%, 11.1%) sit at nearly the same operating point despite sharing no parameters.

**The best method ships the worst default.** Sweeping `fuzz.ratio` gives 0% detection at its shipped cutoff of 95, 2.8% at 70 with no false positives at all, 6.2% at 65 with 0.7%, and 15.3% at 60 with 4.6%. At comparable recall this dominates  $n$ -gram matching: 6.2% detection at 0.7% error against a whitespace 10-gram’s 7.2% at 3.9%, a fivefold difference. The method that performs best in our comparison is the one whose default disables it, and the remedy is a threshold roughly thirty points lower.

## 12 Contamination Is Concentrated in One Slice

Mapping flagged items back to the Dolci source that triggered them shows overlap is not spread evenly across the remix. Under the OpenThoughts3 filter, which flags 138 measurement items, SYNTHETIC-2 supplies 4.9% of our corpus sample and accounts for 72 of them, a tenfold enrichment. Nemotron Code supplies 5.2% and accounts for 26 ( $3.6\times$ ). By contrast OpenThoughts3 Math supplies 34.8% of the corpus and accounts for 70 ( $1.5\times$ ), and Python Algorithms supplies 21.1% and accounts for 3 ( $0.1\times$ ).

The ratio of triggering documents to distinct flagged items is itself diagnostic. Python




Figure 2: Precision/recall frontier, derived without labels. Rings mark each project’s shipped default. Whole-string fuzzy matching dominates both  $n$ -gram families in the low-error region, but the project using it ships a cutoff of 95, at which it flags nothing. Points beyond 20% false positives are omitted; at  $n=5$  both  $n$ -gram families flag essentially the whole benchmark.

Algorithms triggers 278 documents against only 3 items, the signature of shared boilerplate; SYNTHETIC-2 triggers 109 documents against 72 items, the signature of item-level overlap. Filters reporting only a document count cannot distinguish these, and all five of ours do exactly that.

### 13 Recommendations

- Report the procedure, not the word.** “Decontaminated” is uninformative without the tokenizer, window size, threshold, and benchmark list. Two of our five filters cannot be reproduced from their papers alone.
- Publish a temporal control.** Any project can report its false-positive rate against benchmarks postdating its corpus, at no annotation cost.
- Prefer whole-string similarity, tuned.** In our comparison fuzzy matching dominated  $n$ -gram matching at every operating point, provided its threshold is calibrated rather than left at 95.
- Do not compare contamination rates**

across projects. They are not measurements of the same quantity.

### 14 A Note on Our Own Provenance

Every number above was regenerated from committed code before submission rather than carried over from the exploratory scripts that first produced it. Three did not survive: a threshold sweep and the whole of §12 had been computed against an early 200-item benchmark set and then reported alongside the final 502-item one, and two error fractions had been formed by dividing rounded percentages rather than counts. All three are corrected here.

We report this because it is the same failure the paper documents elsewhere: a figure whose provenance drifts from the pipeline that is supposed to produce it. It was caught only by regenerating from a clean checkout, which is the practice we recommend and not one we can claim was unnecessary. Every table and figure in this paper is emitted by a numbered script, and `make paper` reproduces them all.

### 15 Limitations

**We do not measure score inflation.** This work measures which items filters flag, not what removing them does to a reported benchmark number. Establishing that would require retraining under each filter, which we did not do.

**One domain.** All results are on mathematics. The false-positive measurement replicates on a second corpus (§9), but whether the stockphrase mechanism of §10 generalises to code or natural-language benchmarks is untested, though we would expect analogous conventions. The detection-side results rest on Dolci alone.

**The control assumes a known cutoff.** We use a public release date as an upper bound on corpus assembly. A corpus updated after release, or one whose components postdate the release of the whole, would violate this. We test the assumption directly in §fsec:validate and find no control item present, but the test is specific to this corpus and must be repeated for any other.

**Tulu 3 is evaluated on prose-derived parameters.** Its code declares no defaults, so

we used the values its paper states. A different reading would give different results, which is itself part of our argument.

decon is evaluated on its defaults only. We swept its score threshold but did not explore its other parameters, of which there are many.

## 16 Conclusion

Five published decontamination filters, run over one corpus against one benchmark suite, flag between 0% and 27.5% of the same items. The two that find anything are wrong about more than a tenth of what they report, at a rate stable across a  $31\times$  range of corpus size and traceable to the conventional phrasing of the benchmark genre. A decontamination claim is therefore not a property of a dataset but of a procedure, and the procedures in current use do not agree.

We release the harness, the filter reimplementations, and the temporal control sets, so that any project can measure its own false-positive rate without annotating anything.

## References

Jasper Dekoninck, Nikola Jovanović, Tim Gehrunger, K’ari Rögnvaldsson, Ivo Petrov, Chenhao Sun, and Martin Vechev. 2026. *Beyond benchmarks: MathArena as an evaluation platform for mathematics with LLMs*. *Preprint*, arXiv:2605.00674.

Etash Guha, Ryan Marten, Sedrick Keh, et al. 2025. *Openthoughts: Data recipes for reasoning models*. *Preprint*, arXiv:2506.04178.

Nathan Lambert, Jacob Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, Faeze Brahman, Lester James V. Miranda, et al. 2024. *Tülu 3: Pushing frontiers in open language model post-training*. *Preprint*, arXiv:2411.15124.

Dawei Li, Renliang Sun, Yue Huang, Ming Zhong, Bohan Jiang, Jiawei Han, Xiangliang Zhang, and Wei Wang. 2026a. Preference leakage: A contamination problem in LLM-as-a-judge. In *International Conference on Learning Representations (ICLR)*. ArXiv:2502.01534.

Songze Li, Chuokun Xu, Jiaying Wang, Xueluan Gong, Chen Chen, Jirui Zhang, Jun Wang, Kwok-Yan Lam, and Shouling Ji. 2025. *LLMs cannot reliably judge (yet?): A comprehensive assessment on the robustness of LLM-as-a-judge*. *Preprint*, arXiv:2506.09443.

Yanhang Li, Zhichao Fan, and Zexin Zhuang. 2026b. *Auditing the audit: Five failure modes in benchmark-validity audits*. *Preprint*, arXiv:2607.02586.

Niklas Muennighoff, Zitong Yang, Weijia Shi, Xiang Lisa Li, Li Fei-Fei, Hannaneh Hajishirzi, Luke Zettlemoyer, Percy Liang, Emmanuel Candès, and Tatsunori Hashimoto. 2025. *s1: Simple test-time scaling*. *Preprint*, arXiv:2501.19393.

Erfan Nourbakhsh, Mohammad Sadegh Sirjani, Amir Mousavi, Khoa Nguyen, John Quarles, Mimi Xie, and Rocky Slavin. 2026. Are LLM benchmarks already contaminated? a systematic review of contamination detection methods. In *Proceedings of the Workshop on Generation, Evaluation and Metrics (GEM)*.

Yonatan Oren, Nicole Meister, Niladri Chatterji, Faisal Ladhak, and Tatsunori B. Hashimoto. 2024. Proving test set contamination in black box language models. In *International Conference on Learning Representations (ICLR)*. ArXiv:2310.17623.

Oscar Sainz, Jon Ander Campos, Iker García-Ferrero, Julen Etzaniz, Oier Lopez de Lacalle, and Eneko Agirre. 2023. *NLP evaluation in trouble: On the need to measure LLM data contamination for each benchmark*. *Preprint*, arXiv:2310.18018. Findings of EMNLP.

Ibne Farabi Shihab, Sanjeda Akter, and Anuj Sharma. 2026. *When is benchmark contamination detectable? information limits and power-calibrated audits*. *Preprint*, arXiv:2608.07914.

Haoxiang Sun, Yingqian Min, Zhipeng Chen, Wayne Xin Zhao, and Ji-Rong Wen. 2025. *Challenging the boundaries of reasoning: An olympiad-level math benchmark for large language models*. *Preprint*, arXiv:2503.21380. Accepted to ACL 2026 Main.

Team Olmo. 2025. *Olmo 3*. *Preprint*, arXiv:2512.13961. V1 December 2025; v2 April 2026.

Han Wang, Haoyu Li, Brian Ko, and Huan Zhang. 2026. On the fragility of benchmark contamination detection in reasoning models. In *International Conference on Learning Representations (ICLR)*. ArXiv:2510.02386.

Shuo Yang, Wei-Lin Chiang, Lianmin Zheng, Joseph E. Gonzalez, and Ion Stoica. 2023. *Rethinking benchmark and contamination for language models with rephrased samples*. *Preprint*, arXiv:2311.04850.

Wojciech Zarzecki, Jan Dubiński, and Sebastian Cygert. 2026. *The reliability gap in benchmark auditing: Distribution shift and scale as failure modes of contamination detection*. *Preprint*, arXiv:2606.03305.

Zheng Zhang, Qi Liu, Siyuan Liang, Ning Li, Zirui  
Hu, Weibo Gao, Rui Li, Zhenya Huang, and  
Leszek Rutkowski. 2026. Controllable contam-  
ination detection for reliable LLM evaluation  
with statistical guarantees. In *Proceedings of  
the 64th Annual Meeting of the Association for  
Computational Linguistics (ACL)*.